
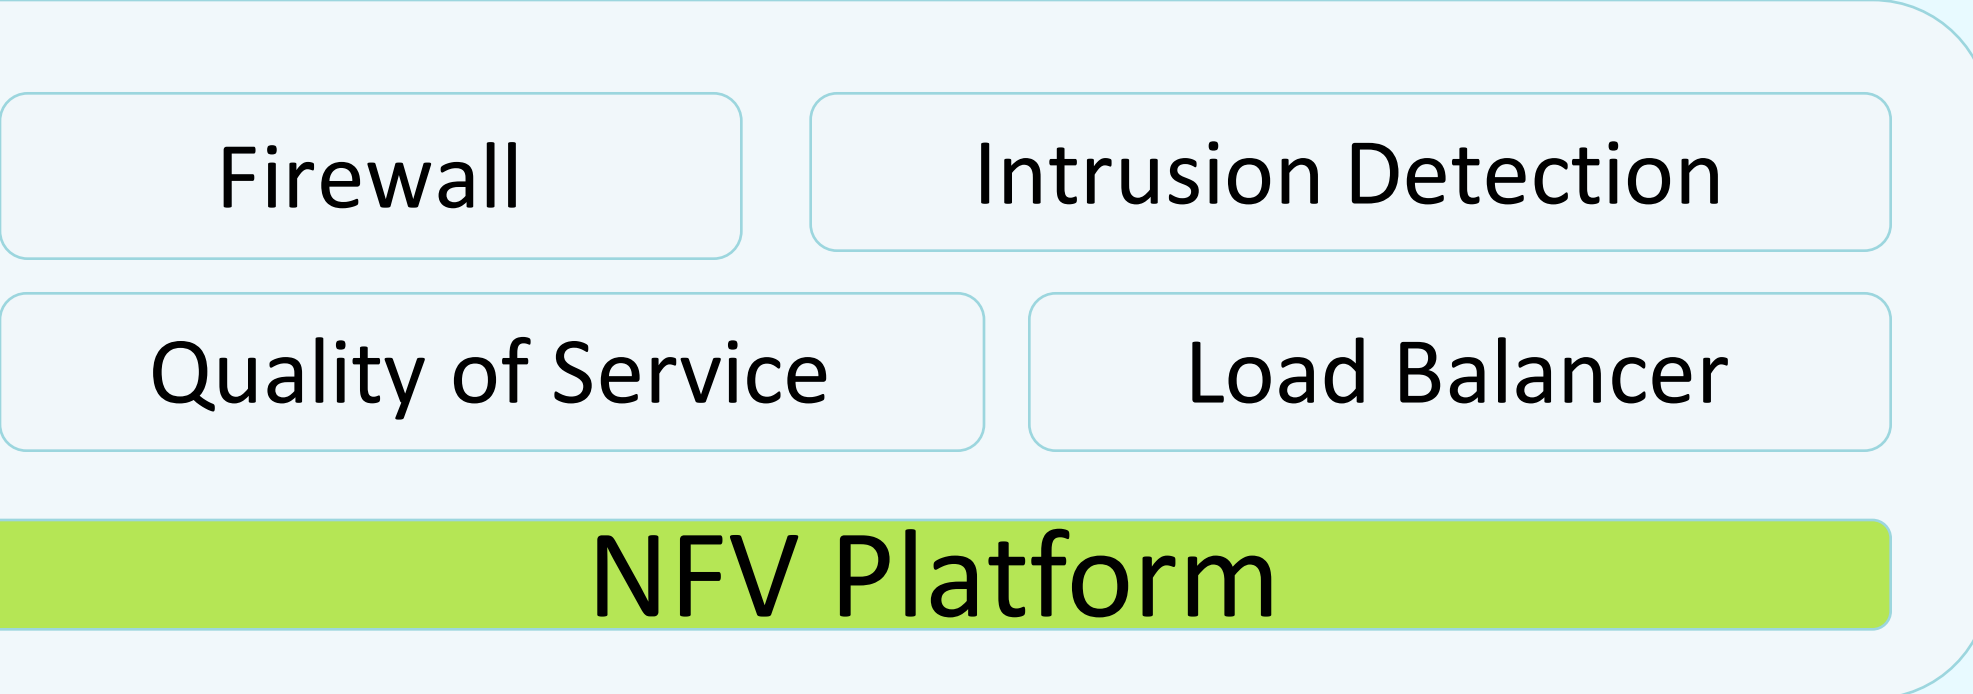


Abstract

- Traditionally, networks are comprised of individual **hardware components** called network functions:



- This model is very **expensive** and **inflexible**
- Trends in networking produced network function virtualization (NFV)
 - Cost effectiveness of software
 - Flexibility of software
 - All network functions on one host

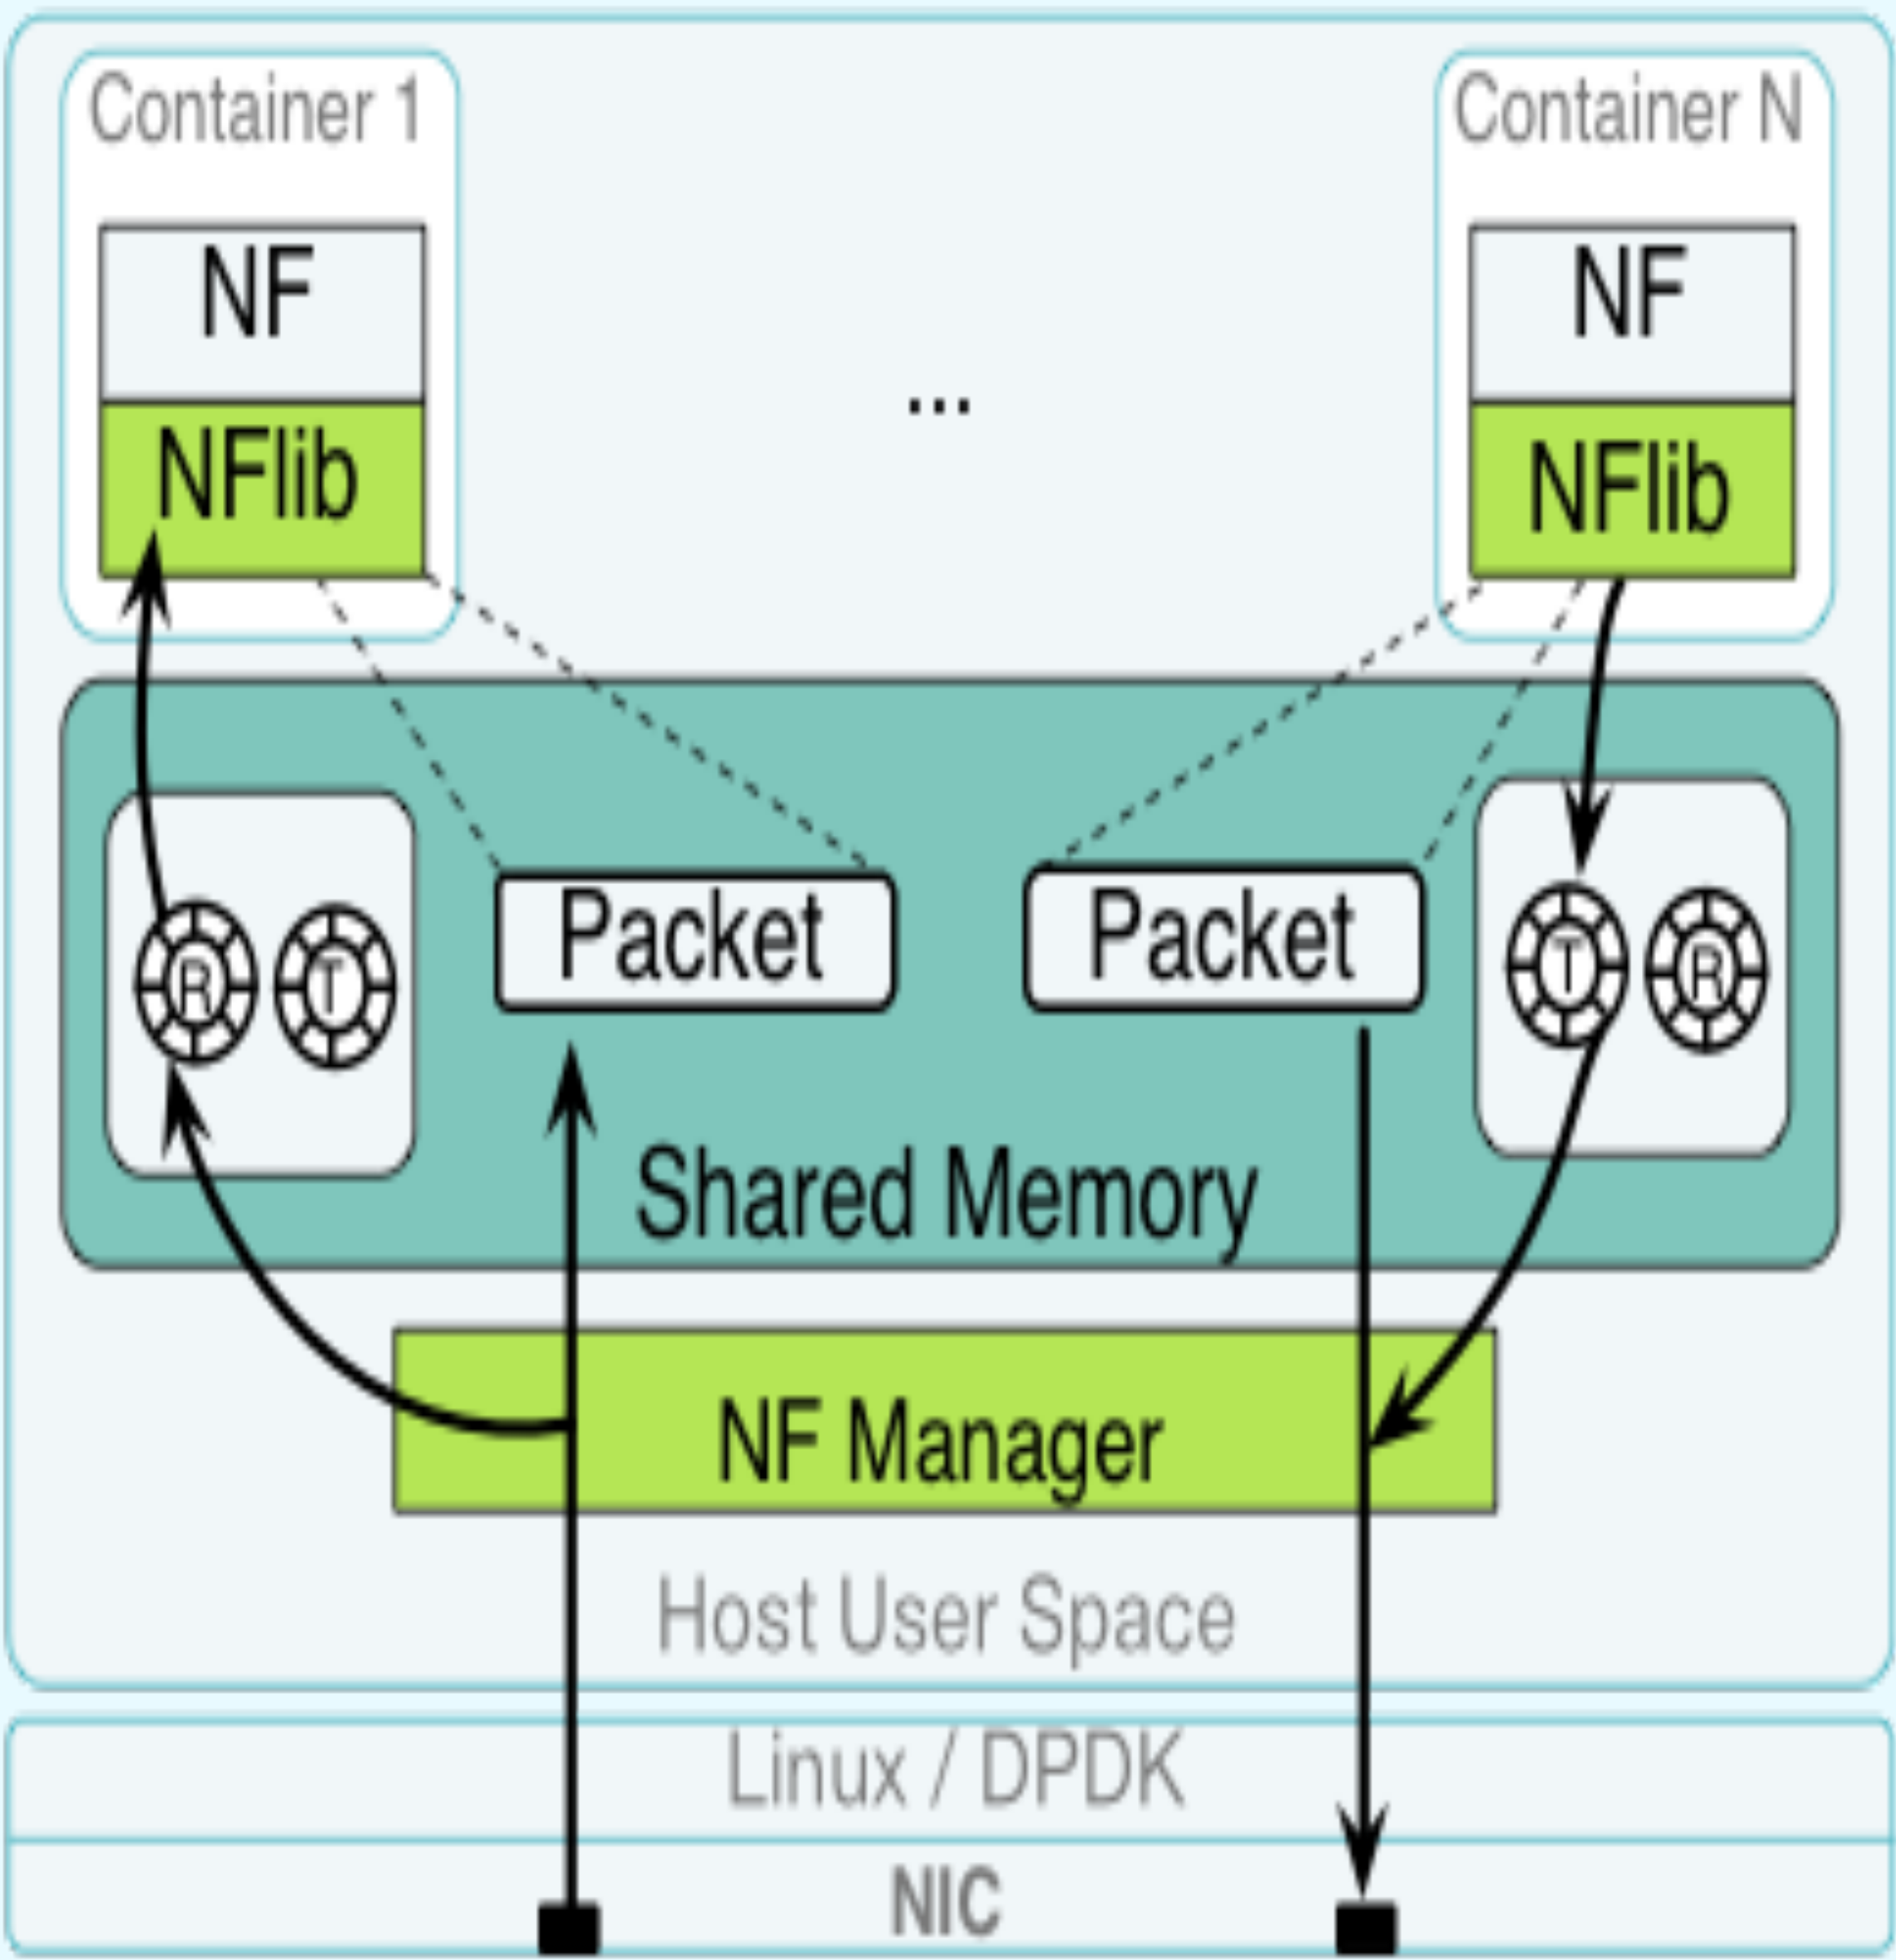


Challenges

- Modern NFV technology still does not provide an elastic framework
 - Adding new NF requires network downtime
- Modern NFV technology is not able to perform at the same line rates as hardware networks

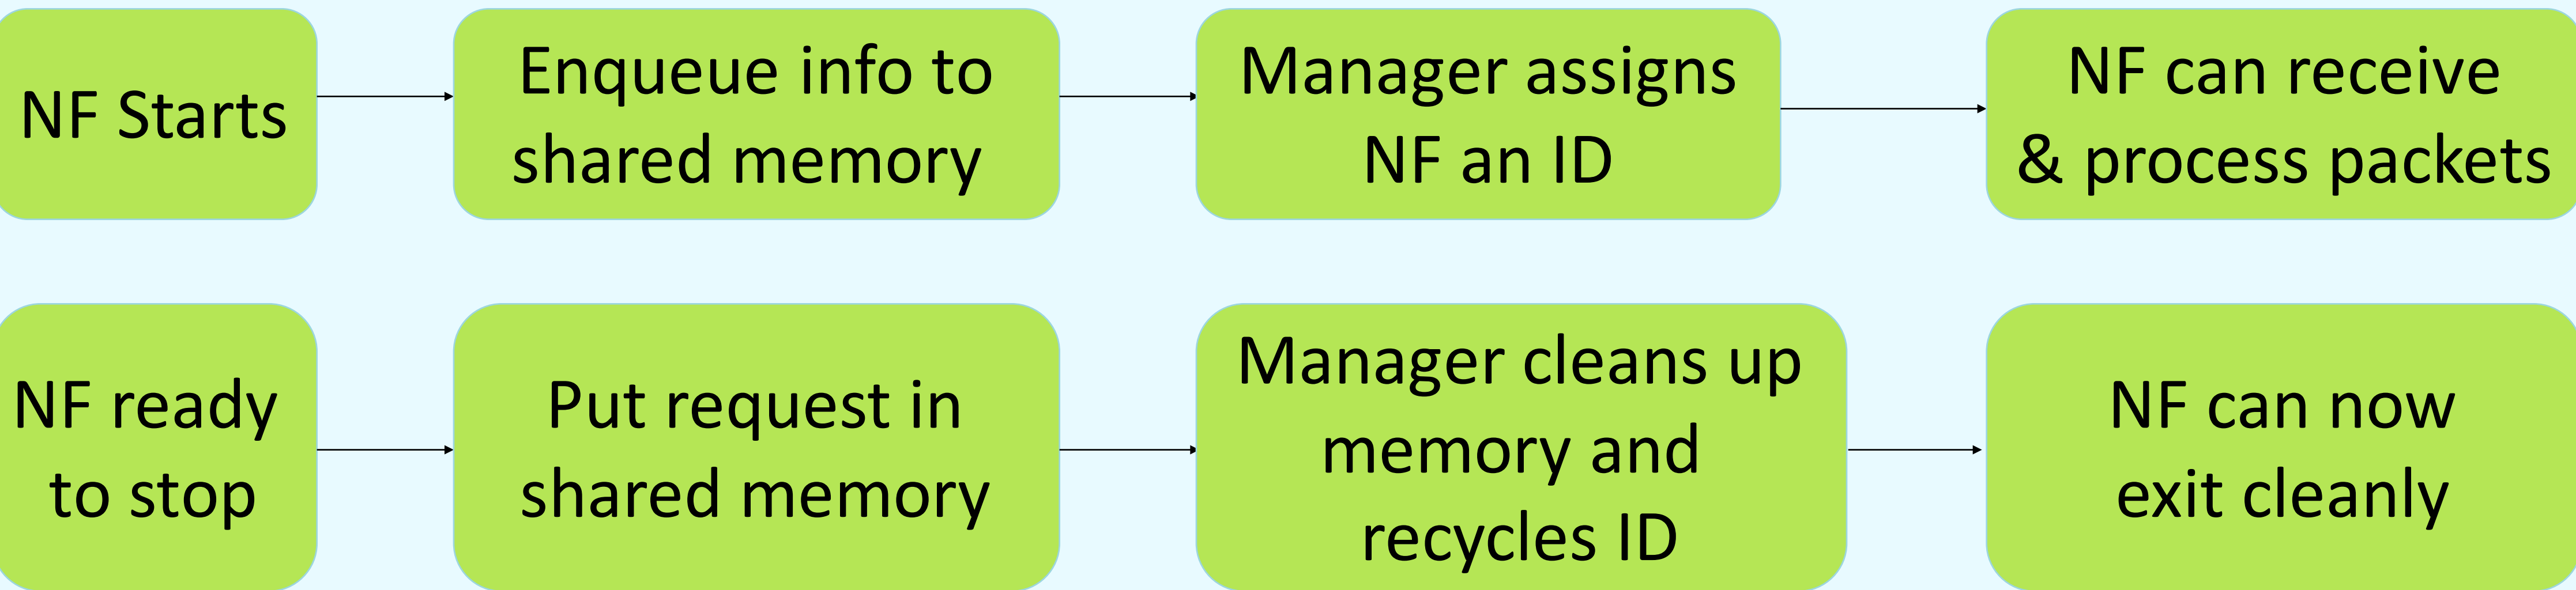
openNetVM

- Container based NFs:** Ease of user-space process management brought to networking
- NF Manager:** Orchestrates traffic flow between various NFs to bring elasticity
- Zero-Copy IO:** Packets DMA'd into shared memory granting NFs direct access to data without copies
- NUMA-Aware:** Maximizes performance by ensuring data in memory is local to a thread's CPU Socket
- Interrupt-Free:** DPDK's poll mode driver allows non-traditional network to process incoming traffic at **10Gbps and beyond**



Dynamic Manager

- As networks grow, more middle boxes need to be deployed to scale efficiently
- openNetVM has a dynamic manager which makes networks elastic
 - Aware of all active and newly created NFs
 - Re organizes data structures upon NF creation and destruction
- Dynamic NF start and stop protocols let the size of the network scale in proportion to traffic without downtime
- System can recover from and restart crashed NFs




TCP/IP Library

- DPDK strips standard packet headers from traffic since it avoids the kernel
- More complicated network functions (IDS, firewall) need TCP/IP packet headers to perform their tasks
- TCP/IP library exposes the standard headers from the packets

TCP/IP Packet																
IP Header	Version		IHL		Type of Service			Total Length								
	Identification					Flags		Fragment Offset								
	Time to Live			Protocol-6 (TCP)			Header Checksum									
	Source Address															
	Destination Address															
	Options					Padding										
	Source Port					Destination Port										
	Sequence Number															
	Acknowledgement Number															
	Data Offset		U		A		P		R		S		F		Window	
			R		S		S		Y		I					
			S		K		H		T		T		N			
Checksum						Urgent Pointer										
TCP Options						Padding										
TCP Data																

Results



- Comparing SR-IOV enabled VMs and DPDK against NetVM (our other system that uses VMs for the same goal), we achieve line rates that are faster than SR-IOV VMs but not faster than raw DPDK
- We expect openNetVM to be as fast as raw DPDK or faster than it since containers are much lighter